

On Sampled Metrics for Item Recommendation

Published in KDD '20, USA

Walid Krichene Steffen Rendle

Google Research
Mountain View, California

Presented by Divyanka Chaudhari, Shreya Gupta

What is Item recommendation?

- Retrieve a **ranked** list of the best items for a user in a particular context.

What is Item recommendation?

- Retrieve a **ranked** list of the best items for a user in a particular context.
- Not all items are necessarily **relevant**.

What is Item recommendation?

- Retrieve a **ranked** list of the best items for a user in a particular context.
- Not all items are necessarily **relevant**.
- Aim to introduce new items to user with a higher probability of them engaging to it.

What is Item recommendation?

- Retrieve a **ranked** list of the best items for a user in a particular context.
- Not all items are necessarily **relevant**.
- Aim to introduce new items to user with a higher probability of them engaging to it.
- Item recommendation algorithms are evaluated by metrics that compare the position of **truly relevant** items among the **recommended items**.

What is Item recommendation?

- Retrieve a **ranked** list of the best items for a user in a particular context.
- Not all items are necessarily **relevant**.
- Aim to introduce new items to user with a higher probability of them engaging to it.
- Item recommendation algorithms are evaluated by metrics that compare the position of **truly relevant** items among the **recommended items**.
- Different methods for comparing an algorithm : AUC, Precision, Recall, Average Precision, NDCG.

Why this paper?

- How does the **sample size** affect output by different metrics?

Why this paper?

- How does the **sample size** affect output by different metrics?
- Are our evaluation metrics **consistent**?

Why this paper?

- How does the **sample size** affect output by different metrics?
- Are our evaluation metrics **consistent**?
- How can we improve the **quality of estimation** when we need to use a sample set?

Evaluating an algorithm

n input items $\rightarrow$ Algorithm A ranks for instance x

Average metric reported

$$\frac{1}{|D|} \sum_{x \in D} M(R(A, x))$$

where $D = \{x_1, x_2, \dots\}$ for multiple instances and $R(A, x)$ are predicted ranks for instance x .

Example: $R(A, x) = \{3, 5\}$

What is Sampling and Consistency?

Sample Set

All relevant + m random irrelevant from total n items.

Consistency

$$\frac{1}{|D|} \sum_{x \in D} M(R(A, x)) > \frac{1}{|D|} \sum_{x \in D} M(R(B, x))$$

$$\iff$$

$$E\left[\frac{1}{|D|} \sum_{x \in D} M(R(A, x))\right] > E\left[\frac{1}{|D|} \sum_{x \in D} M(R(B, x))\right]$$

- If a metric is inconsistent, then measuring M on a subsample is not a good indicator of the true performance of M .

Simplified Metrics

For $|R| = 1$

- $\text{Prec}(r)_k = \delta(r \leq k) \frac{1}{k}$

Simplified Metrics

For $|R| = 1$

- $\text{Prec}(r)_k = \delta(r \leq k)^{\frac{1}{k}}$
- $\text{AP}(r)_k = \delta(r \leq k)^{\frac{1}{r}}$

Simplified Metrics

For $|R| = 1$

- $\text{Prec}(r)_k = \delta(r \leq k) \frac{1}{k}$
- $\text{AP}(r)_k = \delta(r \leq k) \frac{1}{r}$
- $\text{Recall}(r)_k = \delta(r \leq k)$

Simplified Metrics

For $|R| = 1$

- $\text{Prec}(r)_k = \delta(r \leq k) \frac{1}{k}$
- $\text{AP}(r)_k = \delta(r \leq k) \frac{1}{r}$
- $\text{Recall}(r)_k = \delta(r \leq k)$
- $\text{AUC}(r)_k = \frac{n-r}{n-1}$

Simplified Metrics

For $|R| = 1$

- $\text{Prec}(r)_k = \delta(r \leq k) \frac{1}{k}$
- $\text{AP}(r)_k = \delta(r \leq k) \frac{1}{r}$
- $\text{Recall}(r)_k = \delta(r \leq k)$
- $\text{AUC}(r)_k = \frac{n-r}{n-1}$
- $\text{NDCG}(r)_n = \delta(r \leq k) \frac{1}{\log_2(r+1)}$

Simplified Metrics: Example

Example: Recommenders A, B and C, $n = 10000$, $|D| = 5$ instances with 1 relevant item in each instance.

- Recommender A : never good or terrible, ranks the relevant item at position 100 each time.
- Recommender B : ranks relevant item at position 40 for 2 instances.
- Recommender C : ranks relevant item at position 2 in one of the instance, but never able to achieve a good rank for other 4 instances.

Simplified Metrics: Example

	Predicted Ranks	AUC	AP	NDCG	Recall@10
A	100, 100, 100, 100, 100	0.990	0.010	0.150	0.000
B	40, 40, 8437, 9266, 4482	0.555	0.010	0.122	0.000
C	212, 2, 743, 5342, 1548	0.843	0.101	0.208	0.200

Table: Toy Example of evaluating three recommenders A, B, and C on five instances.

- Equally cares about ranks in AUC
- In top heavy metrics C performs better

Simplified Metrics: Example




Figure: Metrics over the whole set of 10000 items.

Simplified Metrics: Example




Figure: Contributions of the top 100 ranks.

Simplified Metrics: Example





Figure: Visualization of metric vs. predicted rank for $n = 10000$. The left side shows the metrics over the whole set of 10000 items. The right side zooms onto the contributions of the top 100 ranks. Side by side comparison.

- All metrics besides AUC are top heavy and almost completely ignore the tail.

Understanding the problem

- **Previous example:** A, B, C recommenders. $n = 10000, |D| = 5$,
 A : 100 each time, B : 40 twice, C : 2 once
- Sampled set: 1 relevant item + $m = 99$ random irrelevant items
- Repeated 1000 times and average and standard deviation is computed.

	Predicted Ranks	AUC	AP	NDCG	Recall@10
A	100, 100, 100, 100, 100	0.990 ± 0.004	0.630 ± 0.129	0.724 ± 0.097	1.000 ± 0.000
B	40, 40, 8437, 9266, 4482	0.555 ± 0.014	0.336 ± 0.073	0.444 ± 0.054	0.400 ± 0.000
C	212, 2, 743, 5342, 1548	0.843 ± 0.014	0.325 ± 0.050	0.460 ± 0.039	0.567 ± 0.092

Figure: Sampled evaluation for the recommenders A, B, C .

Understanding the problem




Figure: Expected Sampling average Precision for the running example.

Understanding the problem




Figure: Expected Sampling average Precision for the running example.

Understanding the problem




Figure: Expected Sampling average Precision for the running example.

Understanding the problem




Figure: Expected Sampling average Precision for the running example.

Understanding the problem





Figure: Expected Sampling NDCG and Recall@10 for the running example.

Understanding the problem




Figure: Expected Sampling AUC for the running example.



Expectation on metrics

- The rank $\tilde{r}$ (measured rank on sample) obtained from the sampling process follows a binomial distribution.



$$\tilde{r} \sim B(m, \frac{r-1}{n-1}) + 1$$

- AUC is linear in r (true rank of unique relevant item). This gives us the expectation of AUC on $\tilde{r}$ as $AUC_n(r)$ which means that measurements created by sampling are **unbiased estimators** of the exact AUC.
- Expectation of Recall and Precision is a **cumulative distribution of binomial distribution**.

Expectation on metrics



Expectation on metrics



Metrics for small sample size

- In this case $\tilde{r} \in \{1, 2\}$ and for any metric M ,
$$E[M(\tilde{r})] = p(\tilde{r} = 1)M(1) + (1 - p(\tilde{r} = 1))M(2)$$

Metrics for small sample size

- In this case $\tilde{r} \in \{1, 2\}$ and for any metric M ,
$$E[M(\tilde{r})] = p(\tilde{r} = 1)M(1) + (1 - p(\tilde{r} = 1))M(2)$$
- This gives a linear expression in r with its coefficient as $M(2) - M(1)$.

Metrics for small sample size

- In this case $\tilde{r} \in \{1, 2\}$ and for any metric M ,
$$E[M(\tilde{r})] = p(\tilde{r} = 1)M(1) + (1 - p(\tilde{r} = 1))M(2)$$
- This gives a linear expression in r with its coefficient as $M(2) - M(1)$.
- For metrics which cannot differentiate between first and second position (Precision, Recall @ $k \geq 2$), M is **constant** hence not useful.

Metrics for small sample size

- In this case $\tilde{r} \in \{1, 2\}$ and for any metric M ,
$$E[M(\tilde{r})] = p(\tilde{r} = 1)M(1) + (1 - p(\tilde{r} = 1))M(2)$$
- This gives a linear expression in r with its coefficient as $M(2) - M(1)$.
- For metrics which cannot differentiate between first and second position (Precision, Recall @ $k \geq 2$), M is **constant** hence not useful.
- For a reasonable metric, the coefficient of r should be **negative**.

Metrics for small sample size



- In this case $\tilde{r} \in \{1, 2\}$ and for any metric M ,
 $E[M(\tilde{r})] = p(\tilde{r} = 1)M(1) + (1 - p(\tilde{r} = 1))M(2)$
- This gives a linear expression in r with its coefficient as $M(2) - M(1)$.
- For metrics which cannot differentiate between first and second position (Precision, Recall @ $k \geq 2$), M is **constant** hence not useful.
- For a reasonable metric, the coefficient of r should be **negative**.
- This shows that any sensible metric gives the **same qualitative** statement for $m = 1$.

Can we design a sampled metric $\hat{M}$, a function from $\{1, \dots, m+1\}$ to $\mathbf{R}$ such that $\hat{M}(\tilde{r})$ provides a good estimate of $M(r)$?

Correction methods:

- Unbiased Estimator of the Rank
- Minimal Bias Estimator
- Bias-Variance Trade-off

Example



- Do recommender algorithms create different ranking distributions?
- Are results from sampled metrics and exact metrics inconsistent?
- Can corrections help to get more reliable results?

Dataset: Binarized Movielens 1M, split, sampling size ($m = 100$) and metrics (Recall@10 and NDCG@10).

Experiments

Figure: Evaluation of three recommenders (X, Y and Z) on the Movielens dataset. Sampled metrics are inconsistent with the exact metrics. Corrected metrics produce the correct relative ordering in expectation.

	Recommender	Exact	Sampled (uncorrected)	Rank Estimate	Sampled with Correction				
					CLS	BV 1	BV 0.1	BV 0.01	BV 0.001
Recall	X	7.60	66.19±0.25	17.46±0.32	8.52±0.16	4.71±0.07	6.49±0.25	7.18±0.59	7.32±1.17
	Y	8.84	56.51±0.22	17.26±0.28	8.42±0.14	4.60±0.07	6.95±0.21	8.18±0.51	8.54±1.07
	Z	9.42	54.20±0.22	18.67±0.32	9.10±0.15	4.97±0.07	7.44±0.24	8.63±0.59	9.08±1.20
NDCG	X	3.76	39.21±0.20	17.46±0.32	3.99±0.07	2.16±0.04	3.00±0.12	3.34±0.32	3.41±0.71
	Y	4.59	34.82±0.16	17.26±0.28	3.94±0.06	2.12±0.03	3.24±0.10	3.85±0.28	4.03±0.66
	Z	4.79	35.34±0.16	18.67±0.32	4.27±0.07	2.29±0.04	3.46±0.12	4.05±0.32	4.31±0.74
AP	X	3.75	32.55±0.21	18.12±0.31	3.58±0.06	2.44±0.03	3.13±0.09	3.37±0.21	3.42±0.49
	Y	4.32	30.01±0.20	17.81±0.28	3.54±0.06	2.32±0.03	3.19±0.07	3.62±0.18	3.73±0.45
	Z	4.44	30.71±0.21	19.20±0.31	3.82±0.06	2.45±0.03	3.38±0.08	3.79±0.21	3.97±0.51
AUC	X	89.13	89.12±0.04	89.24±0.04	89.12±0.04	88.36±0.04	89.04±0.04	89.11±0.04	89.12±0.04
	Y	85.33	85.33±0.04	85.48±0.04	85.32±0.04	84.63±0.04	85.26±0.04	85.32±0.04	85.32±0.04
	Z	74.73	75.04±0.23	75.24±0.20	75.04±0.21	74.51±0.23	75.02±0.20	75.02±0.24	75.02±0.23

Suggestions and Conclusion

Conclusion: Most metrics are inconsistent under sampling and can lead to false discoveries.

Suggestions:

- Sampling should be avoided during evaluation.
- Low-variance can give false confidence in wrong conclusion.
- If not avoidable: corrected metrics can reduce bias.

All the tables and figures have been taken from the presented paper.

Thank You!

divyanka.me/coq301.pdf